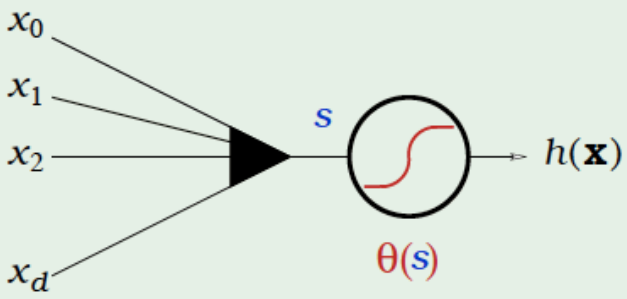
CSE4088 Introduction to Machine Learning

Neural Networks

Slides are adopted from lecture notes of Yaser Abu-Mostafa

Review of last lecture

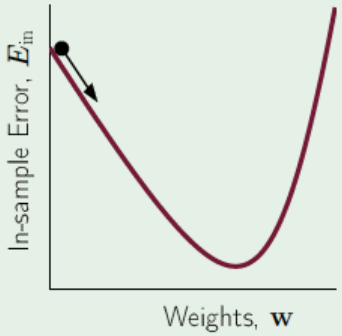
- Logistic regression



- Likelihood measure

$$\prod_{n=1}^{N} P(y_n \mid \mathbf{x}_n) = \prod_{n=1}^{N} \theta(y_n \mathbf{w}^{\mathsf{T}} \mathbf{x}_n)$$

- Gradient descent



- Initialize $\mathbf{w}(0)$
- For $t = 0, 1, 2, \cdots$ [to termination]

$$\mathbf{w}(t+1) = \mathbf{w}(t) - \eta \, \nabla E_{\text{in}}(\mathbf{w}(t))$$

- Return final $\mathbf{w}$

Outline

- Stochastic Gradient Descent
- Neural network model
- Backpropagation model

Stochastic gradient descent

GD minimizes:

$$E_{\text{in}}(\mathbf{w}) = \frac{1}{N}\sum_{n=1}^{N} \underbrace{\mathbf{e}\left(h(\mathbf{x}_n), y_n\right)}_{\ln\left(1+e^{-y_n \mathbf{w}^{\mathsf{T}} \mathbf{x}_n}\right)} \longleftarrow \text{in logistic regression}$$

by iterative steps along $-\nabla E_{\text{in}}$:

$$\Delta \mathbf{w} = -\eta \nabla E_{\text{in}}(\mathbf{w})$$

∇E_{in} is based on all examples $(\mathbf{x}_n, y_n)$

"**batch**" GD

The stochastic aspect

Pick one $(\mathbf{x}_n, y_n)$ at a time. Apply GD to $\mathbf{e}\left(h(\mathbf{x}_n), y_n\right)$

"Average" direction:

$$\mathbb{E}_n\left[-\nabla \mathbf{e}\left(h(\mathbf{x}_n), y_n\right)\right] = \frac{1}{N}\sum_{n=1}^{N} -\nabla \mathbf{e}\left(h(\mathbf{x}_n), y_n\right)$$

$$= -\nabla E_{\text{in}}$$

randomized version of GD

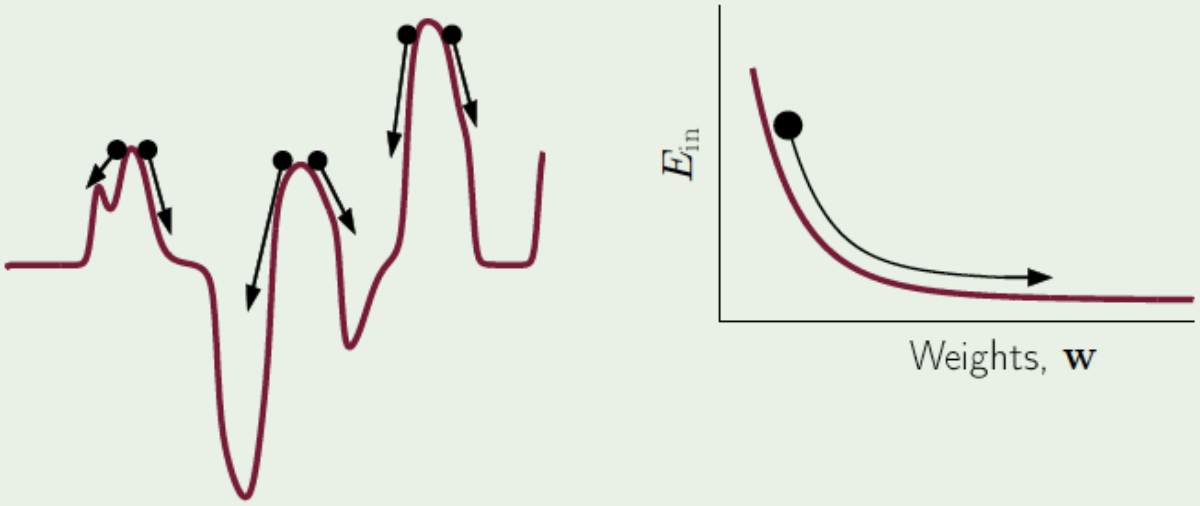
stochastic gradient descent (SGD)

Benefits of SGD

1. cheaper computation
2. randomization
3. simple

Rule of thumb:

$\eta = 0.1$ works



randomization helps

Randomization helps to:

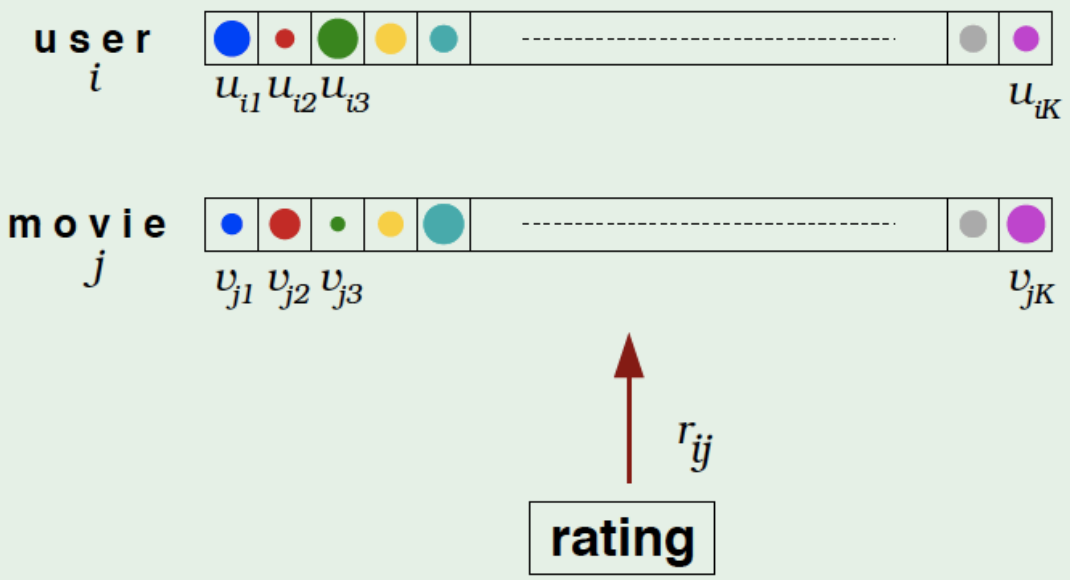
1) Escape shallow minima (e.g. Leftmost minima above)
2) Avoid early termination in flat error regions (right figure)

SGD in action

- Movie ratings

$$e_{ij} = \left(r_{ij} - \sum_{k=1}^{K} u_{ik} v_{jk} \right)^2$$

(To minimize the error we take one sample at a time. The factors are the parameters we are trying to find that minimizes the errors. There are 2K parameters. Take the partial derivative of e with respect to every parameter and change every factor a little bit towards creating the rating we want in proportion to the gradient.)



The user will be described by a number of factors: does he like comedy, action etc.?
We will also describe each movie by the same factors. Is the movie about comedy, action etc.?
Match the user and movie by taking the inner product of their vectors.

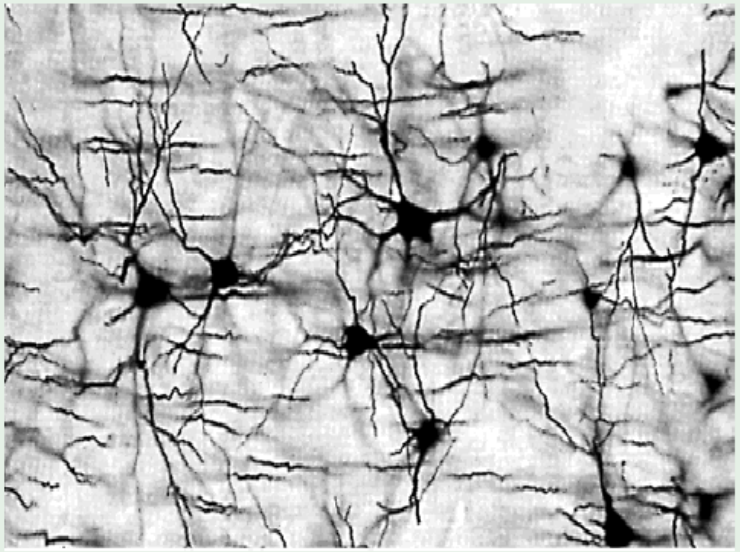
Outline

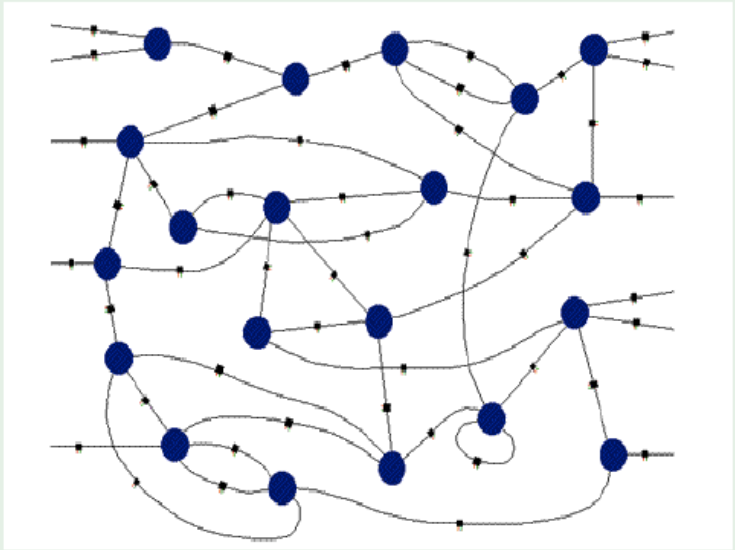
- Stochastic Gradient Descent
- Neural network model
- Backpropagation model

Biological inspiration

- We want to replicate the biological function of learning

biological function ⟶ biological structure





Biological inspiration

biological function ⟶ biological structure

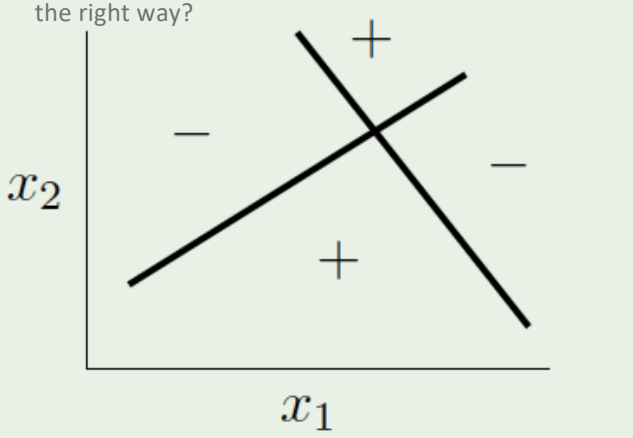


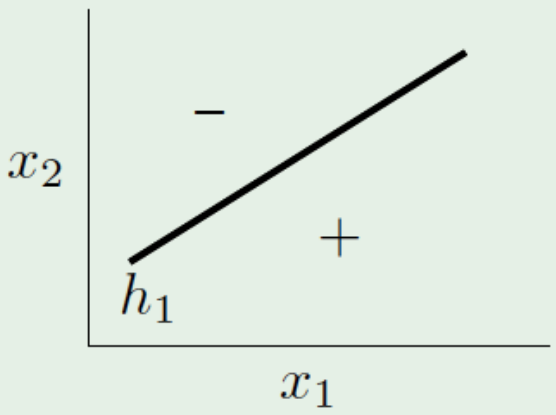


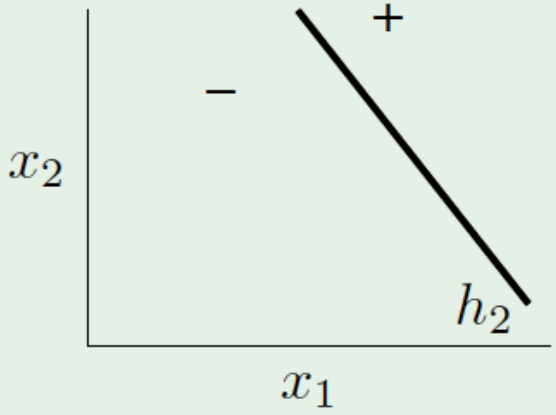
Combining perceptrons

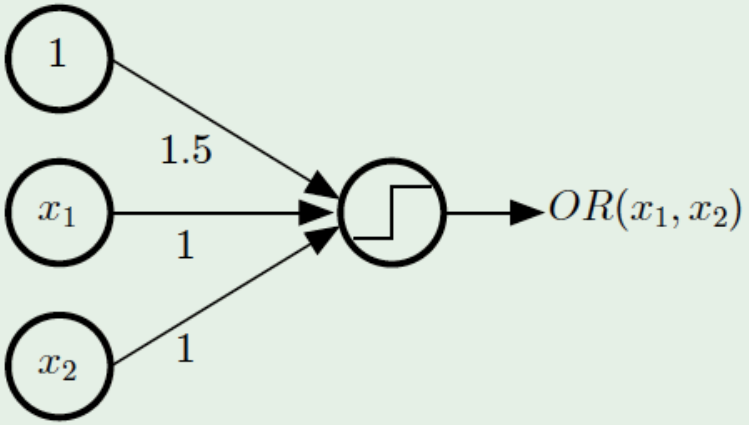
A single perceptron can not shatter 4 points.
Can we do it with more than one perceptrons arranged in the right way?

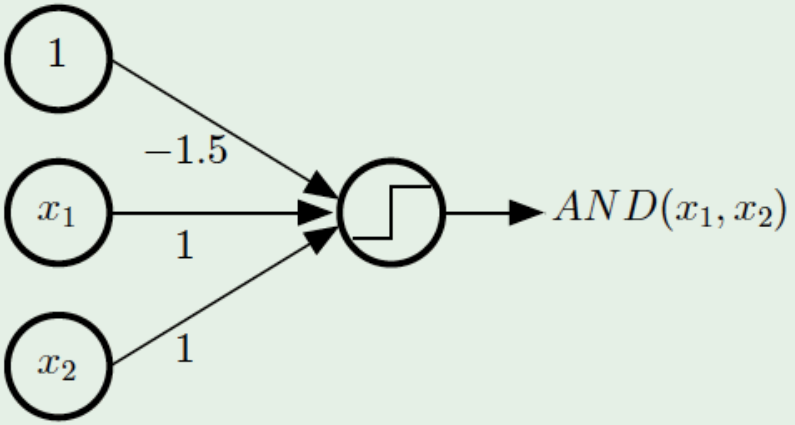
We can XOR h_1 and h_2 to get the classifier on the left



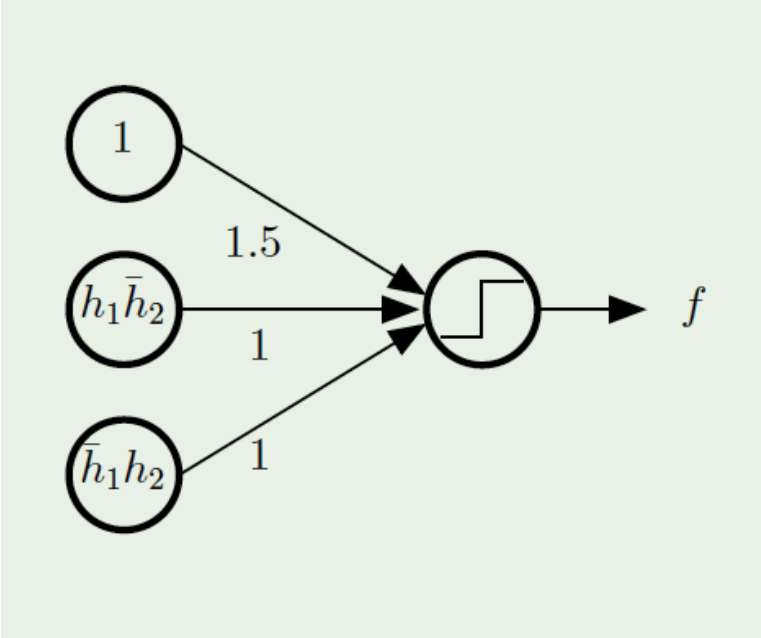




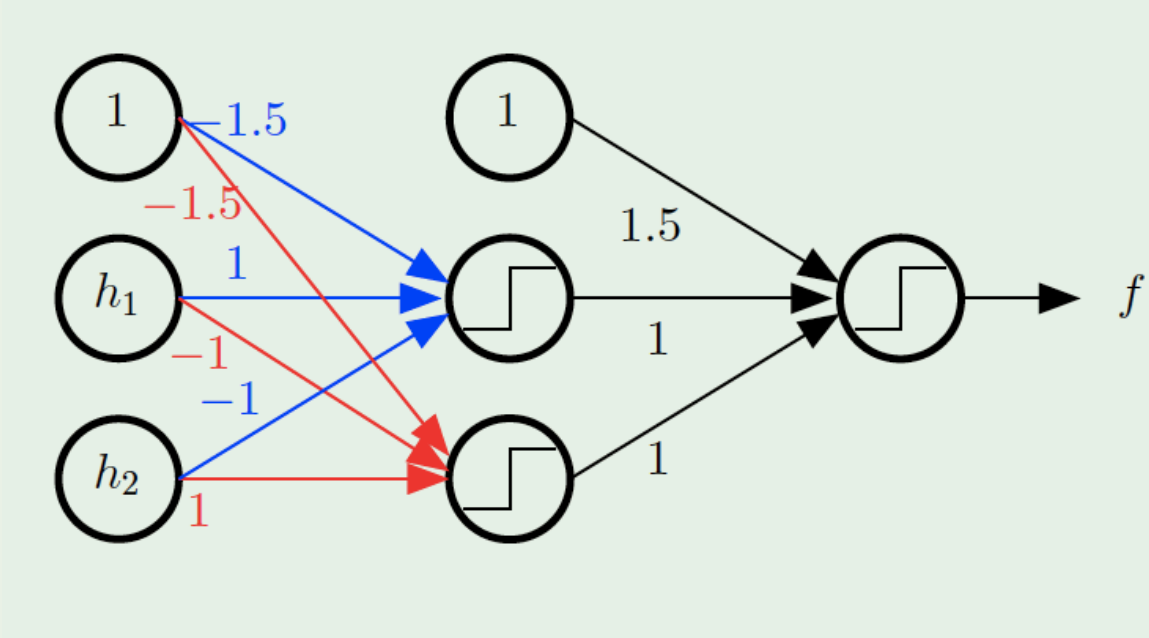




Creating layers



XOR(h_1,h_2)=$OR(h_1\overline{h_2},\ \overline{h_1}h_2)$

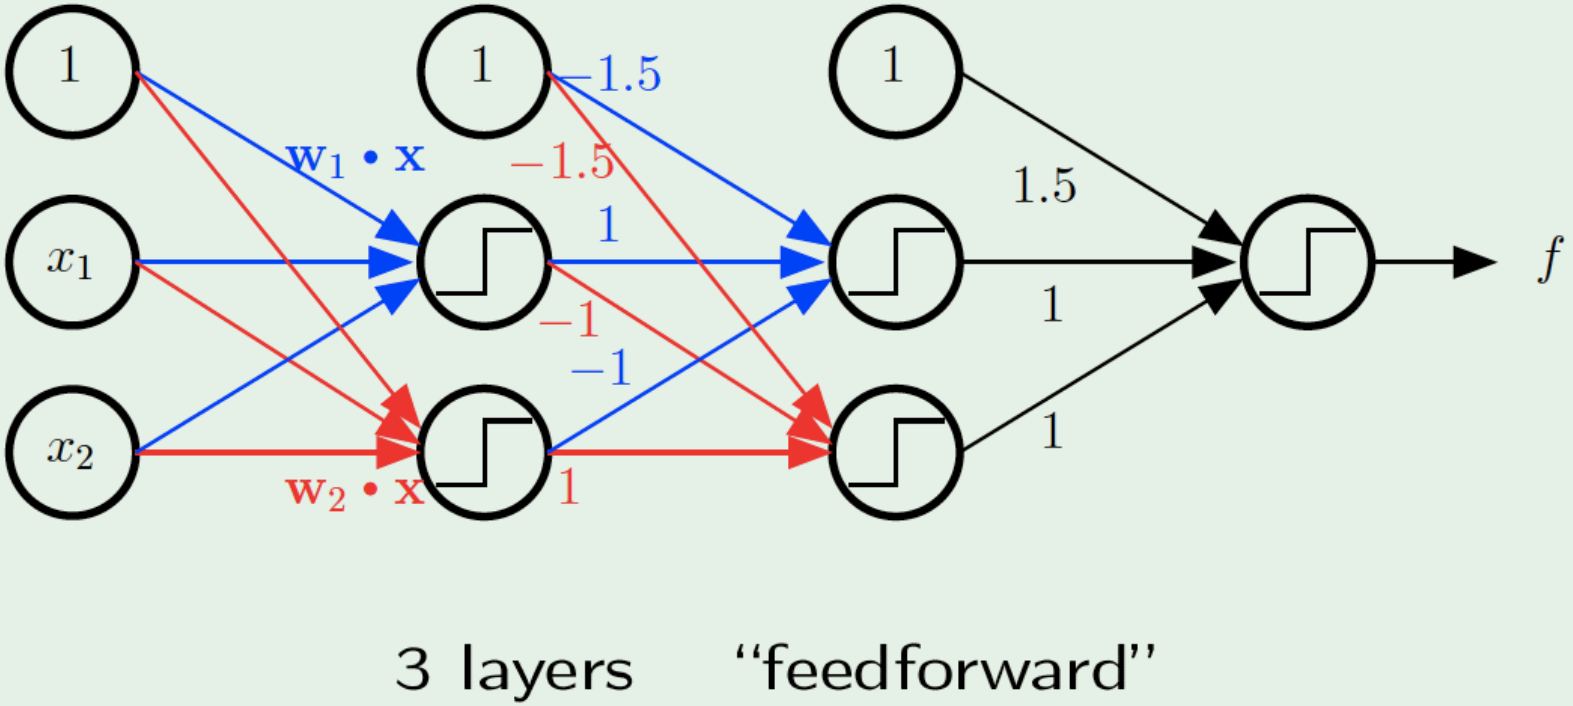


Blue lines implement $h_1\overline{h_2}$

Red lines implement $\overline{h_1}h_2$

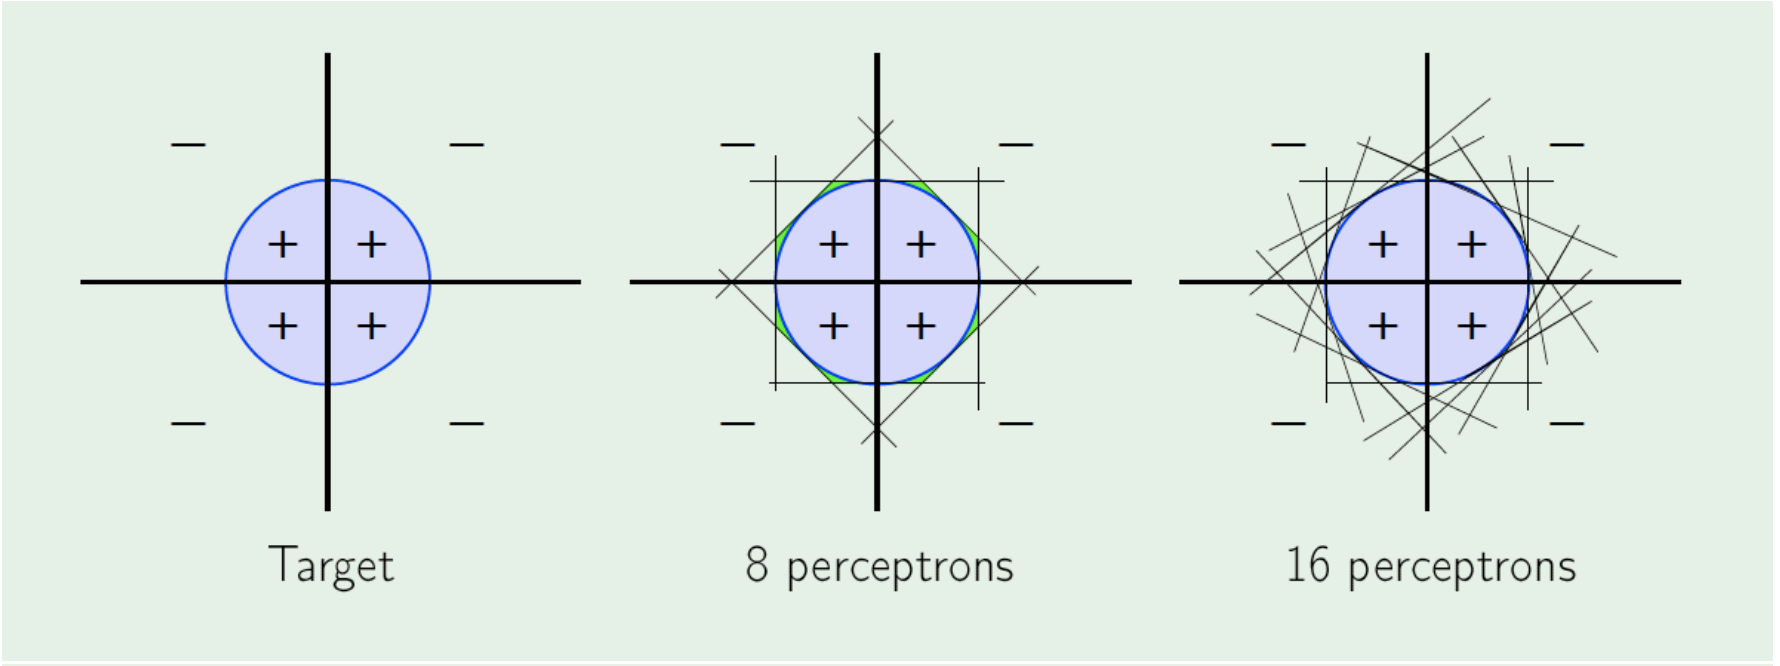
XOR(h_1,h_2)=$OR(h_1\overline{h_2},\ \overline{h_1}h_2)$

The multilayer perceptron



Replace h_1 and h_2 with perceptrons that use the inputs .

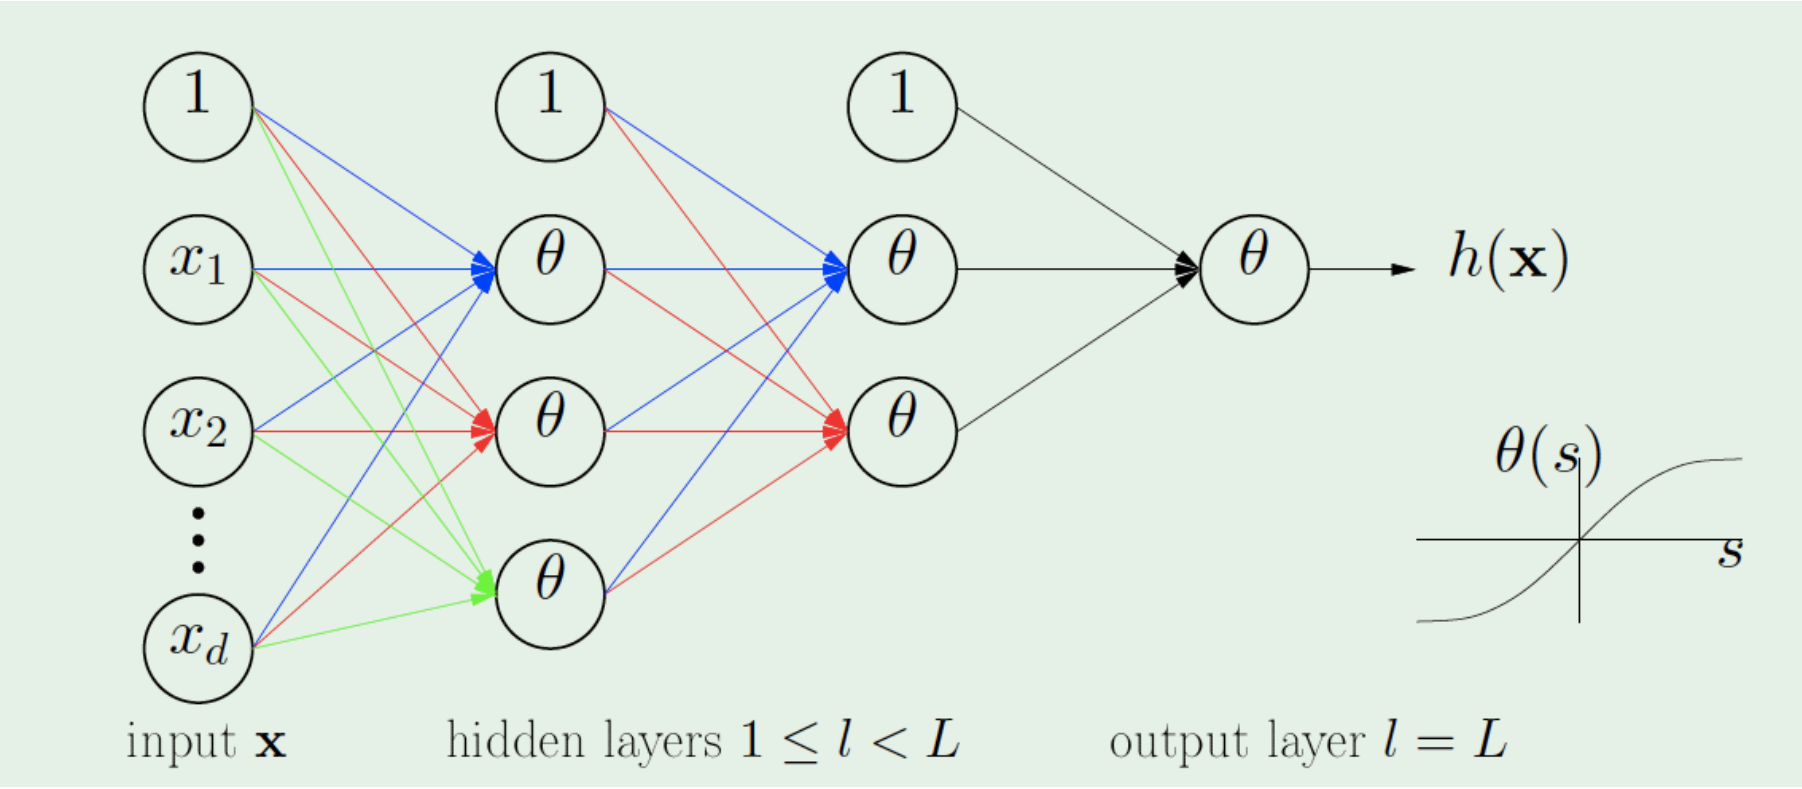
A powerful model



A sufficient number of neurons **can approximate any function**. Two red flags:

- **Generalization**: Since VC dimension increases as the number of neurons increase, we must provide more data.
- **Optimization**: It becomes a difficult combinatorial optimization problem. Solution: we can use a soft threshold and differentiate to optimize (we can use the gradient descent!).

The neural network

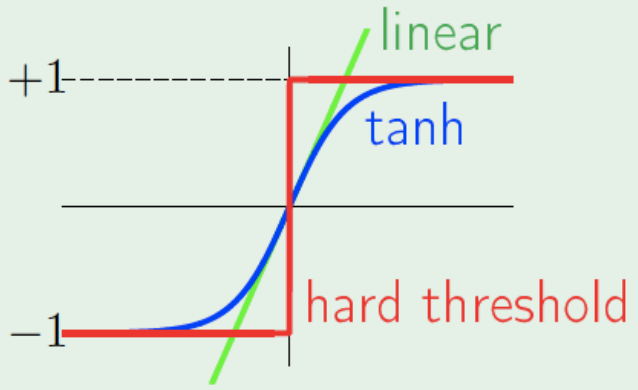


How the network operates

$$w_{ij}^{(l)} \quad \begin{cases} 1 \le l \le L & \text{layers} \\ 0 \le i \le d^{(l-1)} & \text{inputs} \\ 1 \le j \le d^{(l)} & \text{outputs} \end{cases}$$

$$x_j^{(l)} = \theta(s_j^{(l)}) = \theta\left(\sum_{i=0}^{d^{(l-1)}} w_{ij}^{(l)}\, x_i^{(l-1)}\right)$$

Apply $\mathbf{x}$ to $x_1^{(0)} \cdots x_{d^{(0)}}^{(0)} \rightarrow\rightarrow x_1^{(L)} = h(\mathbf{x})$



$$\theta(s) = \tanh(s) = \frac{e^s - e^{-s}}{e^s + e^{-s}}$$

Given the input x, apply the recursive definition through the layers until you get the single output h(x)

Outline

- Stochastic Gradient Descent
- Neural network model
- Backpropagation model (to learn the weights)

Applying SGD

All the weights $\mathbf{w} = \{w_{ij}^{(l)}\}$ determine $h(\mathbf{x})$

Error on example $(\mathbf{x}_n, y_n)$ is

$$\mathbf{e}\left(h(\mathbf{x}_n), y_n\right) = \mathbf{e}(\mathbf{w})$$

To implement SGD, we need the gradient

$$\nabla \mathbf{e}(\mathbf{w}): \quad \frac{\partial\, \mathbf{e}(\mathbf{w})}{\partial\, w_{ij}^{(l)}} \quad \text{for all} \quad i, j, l$$

Take one example, and adjust all the weights of the network in the direction of the negative gradient, according to that single example.

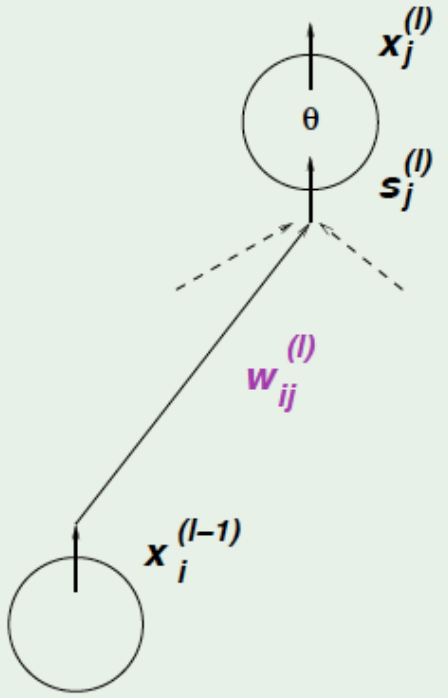
Computing $\frac{\partial\ \mathbf{e}(\mathbf{w})}{\partial\ w_{ij}^{(l)}}$

We can evaluate $\frac{\partial\ \mathbf{e}(\mathbf{w})}{\partial\ w_{ij}^{(l)}}$ one by one: analytically or numerically

A trick for efficient computation:

$$\frac{\partial\ \mathbf{e}(\mathbf{w})}{\partial\ w_{ij}^{(l)}} = \frac{\partial\ \mathbf{e}(\mathbf{w})}{\partial\ s_j^{(l)}} \times \frac{\partial\ s_j^{(l)}}{\partial\ w_{ij}^{(l)}}$$

We have $\frac{\partial\ s_j^{(l)}}{\partial\ w_{ij}^{(l)}} = x_i^{(l-1)}$ We only need: $\frac{\partial\ \mathbf{e}(\mathbf{w})}{\partial\ s_j^{(l)}} = \delta_j^{(l)}$



δ for the final layer

$$\delta_j^{(l)} = \frac{\partial\, \mathbf{e}(\mathbf{w})}{\partial\, s_j^{(l)}}$$

For the final layer $l = L$ and $j = 1$:

$$\delta_1^{(L)} = \frac{\partial\, \mathbf{e}(\mathbf{w})}{\partial\, s_1^{(L)}}$$

Output $h(x_n)$

Target

$$\mathbf{e}(\mathbf{w}) = (\, x_1^{(L)} - y_n)^2$$

$$x_1^{(L)} = \theta(s_1^{(L)})$$

$$\theta'(s) = 1 - \theta^2(s) \quad \text{for the tanh}$$

$$\delta_1^{(L)} = \frac{\partial e(w)}{\partial s_1^{(L)}} = \frac{\partial e(w)}{\partial x_1^{(L)}} \times \frac{\partial x_1^{(L)}}{\partial s_1^{(L)}}$$

$$= 2\left(x_1^{(L)} - y_n\right) \times \theta'\left(s_1^{(L)}\right)$$

$$= 2\left(x_1^{(L)} - y_n\right)\left(1 - \theta^2\left(s_1^{(L)}\right)\right)$$

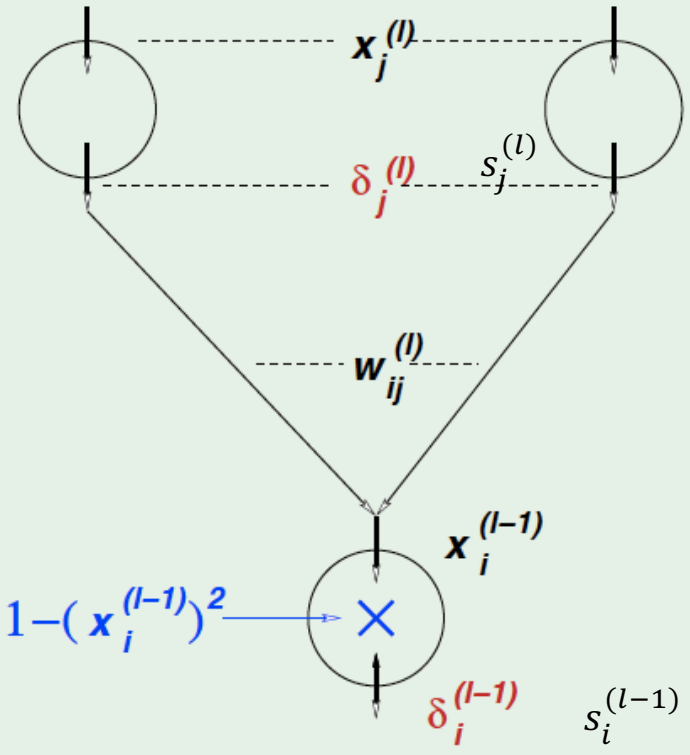
Back propagation of δ

$$\delta_i^{(l-1)} = \frac{\partial\, \mathbf{e}(\mathbf{w})}{\partial\, s_i^{(l-1)}}$$

$$= \sum_{j=1}^{d^{(l)}} \frac{\partial\, \mathbf{e}(\mathbf{w})}{\partial\, s_j^{(l)}} \times \frac{\partial\, s_j^{(l)}}{\partial\, x_i^{(l-1)}} \times \frac{\partial\, x_i^{(l-1)}}{\partial\, s_i^{(l-1)}}$$

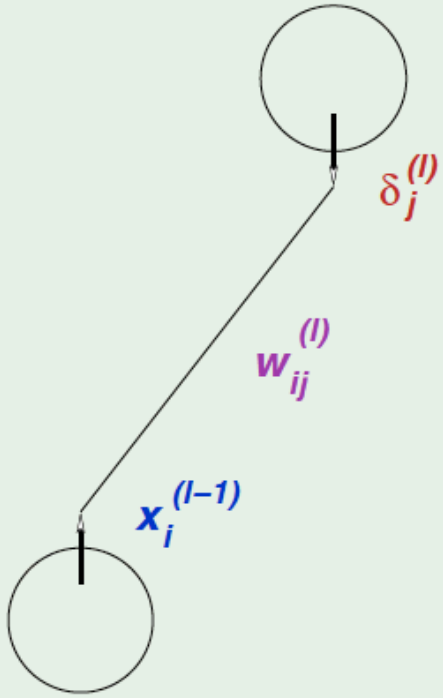
$$= \sum_{j=1}^{d^{(l)}} \delta_j^{(l)} \times w_{ij}^{(l)} \times \theta'(s_i^{(l-1)})$$

$$\delta_i^{(l-1)} = (1 - (x_i^{(l-1)})^2) \sum_{j=1}^{d^{(l)}} w_{ij}^{(l)}\, \delta_j^{(l)}$$



Backpropagation algorithm

1: Initialize all weights $w_{ij}^{(l)}$ at random
2: **for** $t = 0, 1, 2, \ldots$ **do**
3: Pick $n \in \{1, 2, \cdots, N\}$
4: *Forward:* Compute all $x_j^{(l)}$
5: *Backward:* Compute all $\delta_j^{(l)}$
6: Update the weights: $w_{ij}^{(l)} \leftarrow w_{ij}^{(l)} - \eta\, x_i^{(l-1)} \delta_j^{(l)}$
7: Iterate to the next step until it is time to stop
8: Return the final weights $w_{ij}^{(l)}$

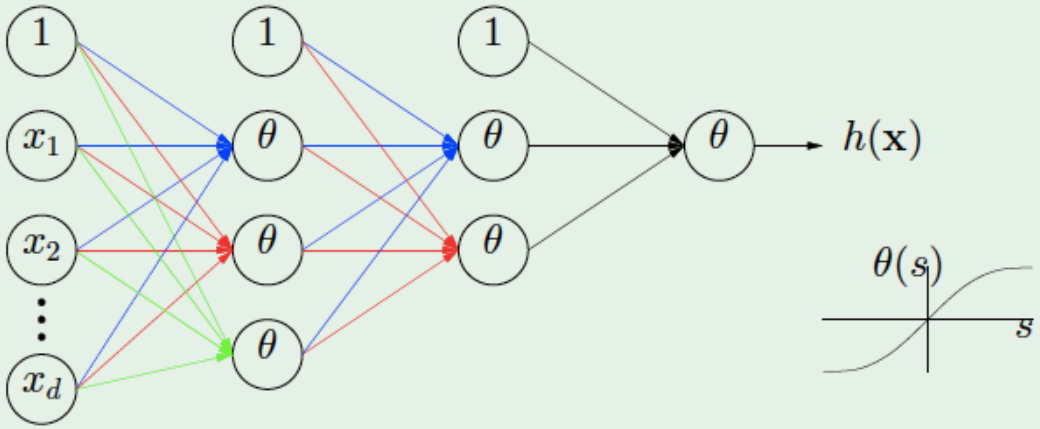


Note: If we initialize the weights to zero, we will be not be able to update the weights.

Final remark: hidden layers

learned nonlinear transform

interpretation?



Note: Hidden layers are performing nonlinear transformations on the outputs of the previous layer. Hence, they are like non-linear features (of features). But they are LEARNED in the neural network.

-- Can we interpret the outputs of the hidden layers? --> Not very easy.

-- The goal of a neural network is to produce the right hypothesis. Not to explain to a human what the hidden layers are doing.